# Streamlined implementation of a machine learning model to classify screen relevant neoplasia using reference-based OTU clustering

Running title: Reference-based OTU clustering for machine learning classification

Courtney R. Armour[1], Kelly L. Sovacool[2], William L. Close[1,*], Begüm D. Topçuoğlu[1,#], Jenna Wiens[3], Patrick D. Schloss [1,†]

[1] Department of Microbiology and Immunology, University of Michigan, Ann Arbor, Michigan, USA

[2] Department of Computational Medicine and Bioinformatics, University of Michigan, Ann Arbor, Michigan, USA

[3] Department of Electrical Engineering and Computer Science, University of Michigan, Ann Arbor, Michigan, USA

[*] Current Affiliation: Bio-Rad Laboratories, Hercules, California, USA

[#] Current Affiliation: Bristol Myers Squibb, Summit, New Jersey, USA

[†] To whom correspondence should be addressed: pschloss@umich.edu

**observation format** (max 1200 words, 2 figures, 25 ref)

## Abstract

Machine learning classification of disease based on the gut microbiome often relies on clustering 16S rRNA gene sequences into operational taxonomic units (OTUs) to quantify microbial composition. The abundance of each OTU is then used to train a classification model. The standard *de novo* approach to clustering sequences into OTUs leverages the similarity of the sequences to each other rather than to a reference database. However, such an approach depends on the sequences in the dataset and therefore OTU assignments can change if new data are added. This lack of stability complicates classification because in order to use the model to classify additional samples, the new sequences must be reclustered with the old data and the model must be retrained with the new OTU assignments. The new reference-based clustering algorithm, OptiFit, addresses this issue by fitting new sequences into existing OTUs. While OptiFit can produce high quality OTU clusters, it is unclear whether this method for fitting new sequence data into existing OTUs will impact the performance of classification models. We used OptiFit to cluster additional data into existing OTU clusters and then evaluated model performance in classifying a dataset containing samples from patients with and without colonic screen relevant neoplasia (SRN). We compared the performance of this model to the standard procedure of *de novo* clustering all the data together. We found that both approaches performed equally well in classifying SRNs. Moving forward, when OTUs are used in classification problems, OptiFit can streamline the process of classifying new samples by avoiding the need to retrain models using reclustered sequences.

## Importance

There is great potential for using microbiome data to non-invasively diagnose people. One of the challenges with using classification models based on the relative abundance of operational taxonomic units (OTUs) is that 16S rRNA gene sequences are assigned to OTUs based on their similarity to other sequences in a dataset. If data are generated from new patients seeking a diagnosis, then a standard approach requires reassigning sequences to OTUs and retraining the classification model. Yet there is a desire to have a single, validated model that can be widely deployed. To overcome this obstacle, we applied the OptiFit clustering algorithm which fits new sequence data to existing OTUs and allows for the reuse of models. A random forest machine learning model implemented using OptiFit performed as well as the traditional reassignment and retrain approach. This result indicates that there is potential for developing and applying machine learning models based on OTU relative abundance data that do not require retraining.

2

30 Gut community composition is useful as a resource for machine learning classification of diseases, such as

31 colorectal cancer (1, 2). Taxonomic composition of microbial communities can be assessed using amplicon

32 sequencing of the 16S rRNA gene, which is the input to classification models. Analysis of 16S rRNA

33 gene sequence data generally relies on similarity-based clustering of sequences into operational taxonomic

34 units (OTUs). The process of OTU clustering can either be reference-based or *de novo*. The quality of

35 OTUs generated with reference-based clustering is generally poor compared to those generated with *de*

36 *novo* clustering (3). While *de novo* clustering produces high-quality OTU clusters where sequences are

37 accurately grouped based on similarity thresholds, the resulting OTU clusters depend on the sequences

38 within the dataset and the addition of new data has the potential to redefine OTU cluster composition.

39 The unstable nature of *de novo* OTU clustering complicates deployment of machine learning models since

40 integration of additional data requires reclustering all the data and retraining the model. The ability to

41 integrate new data into a validated model without reclustering and retraining could allow for the application of

42 a single model that can continually classify new data. Recently, Sovacool *et al.* introduced OptiFit, a method

43 for fitting new sequence data into existing OTUs (4). While OptiFit can effectively fit new sequence data to

44 existing OTU clusters, it is unknown if the use of OptiFit will have an impact on classification performance.

45 Here, we tested the ability of OptiFit to cluster new sequence data into existing OTU clusters for the purpose

46 of classifying disease based on gut microbiome composition.

47 We compared two approaches, one using all the data to generate OTU clusters, and the other generating

48 *de novo* OTU clusters with a portion of the data and then fitting the remaining sequence data to the existing

49 OTUs using OptiFit. In the first approach, all the 16S rRNA sequence data was clustered into OTUs with the

50 OptiClust algorithm in mothur (5). The resulting abundance data was then split into training and testing sets,

51 where the training set was used to tune hyperparameters and ultimately train and select the model. The

52 model was applied to the testing set and performance evaluated (Figure 1A). However, this methodology

53 requires one to regenerate the OTU clusters and retrain the model to classify additional samples. The

54 OptiFit algorithm (4) addresses this problem by enabling new sequences to be clustered into existing OTUs.

55 The OptiFit workflow is similar to the OptiClust workflow, where the data was clustered into OTUs and used

56 to tune hyperparameters and ultimately train the model. Then, we used OptiFit to fit sequence data of

57 samples not part of the original dataset into the existing OTUs, and used the same model to classify the

58 samples (Figure 1B). To test how the model performance compared between these two approaches, we

59 used a publicly available dataset of 16S rRNA gene sequences from stool samples of healthy subjects (n

60 = 261) as well as subjects with SRN consisting of advanced adenoma and carcinoma (n = 229) (1). The

61 dataset was randomly split into an 80% train set and 20% test set. For the standard OptiClust workflow,

the training and test sets were *de novo* clustered together into OTUs, then the resulting abundance table was split into the training and testing set. For the OptiFit workflow, the train set was clustered *de novo* into OTUs, and the remaining test set was fit to the OTU clusters using the OptiFit algorithm. For both workflows, the abundance table of the train set was used to tune hyperparameters and train a random forest model to classify SRN. The test set was classified as either control or SRN using the trained models. To account for variation, the dataset was randomly split 100 times and the process repeated for each of the 100 data splits. By comparing the model performance of classifying the samples in the test dataset between the OptiFit and OptiClust algorithms, we quantified the impact of using OptiFit on model classification performance.

We first examined the quality of the resulting OTU clusters from the two algorithms using the Matthews correlation coefficient (MCC). The MCC score was quantified by examining all pairs of sequences and assessing whether they belonged together in an OTU based on their similarity (5). MCC scores range between negative one and one. A score of negative one means none of the sequences in an OTU are within the similarity threshold and any sequences within the similarity threshold are not in an OTU together. An MCC score of zero means the sequences are randomly clustered. An MCC score of 1 means all sequences in an OTU are within the similarity threshold and all sequence pairs within the similarity threshold are in the same OTU. To ensure that OptiFit is appropriately integrating new sequence data into the existing OTUs, we expected the MCC scores produced by the OptiClust and OptiFit workflows to be similar. Since the data was only clustered once in the OptiClust workflow there was only one MCC score (MCC = 0.884) while the OptiFit workflow produced an MCC score for the OTU clusters from each data split (average MCC = 0.879, standard deviation = 0.002). As expected, the MCC scores were similar between the two workflows. Another metric we examined for the OptiFit workflow was the fraction of sequences from the test set that mapped to the reference OTUs. Any sequences that did not map to reference OTUs were eliminated therefore, if a high percentage of reads did not map we might expect this loss of data to negatively impact classification performance. We found that loss of data was not an issue since on average 99.9% (standard deviation = 0.004) of sequences in the test set mapped to the reference OTUs. These results indicate that OptiFit performed as well as OptiClust when integrating new sequences into the existing OTUs.

After verifying that the quality of the OTUs was consistent between OptiClust and OptiFit, we examined the model performance for classifying samples in the held out test dataset. To evaluate model performance, we used the OTU relative abundances from the training data from the OptiClust and OptiFit workflows to train a model to predict SRNs. Using the predicted and actual diagnosis classification, we calculated the area under the receiver operating characteristic curve (AUROC) for each data split to quantify model performance. During cross-validation (CV) training, the model performance was equivalent between the

4

OptiClust and OptiFit approaches (p-value = 0.13; Figure 2A). The trained model was then applied to the test data classifying samples as either control or SRN. The performance on the test data was equivalent between the OptiClust and OptiFit approaches (p-value = 0.38; Figures 2B and 2C) indicating that new data could be fit to existing OTU clusters using OptiFit without impacting model performance.

We tested the ability of OptiFit to integrate new data into existing OTUs for the purpose of machine learning classification using OTU relative abundance. A potential problem with using OptiFit is that any sequences from the new samples that do not map to the existing OTU clusters will be discarded, resulting in a possible loss of information. However, we demonstrated OptiFit can be used to fit new sequence data into existing OTU clusters and performs equally well in predicting SRN compared to *de novo* clustering all the sequence data together. The ability to integrate data from new samples into existing OTUs enables the implementation of a single machine learning model. This is important for model implementation because not all of the data needs to be available or known at the time of model generation. Since results may depend on the amount of data in the training set, further analysis is needed to determine the number of samples that are necessary in training to build a robust model capable of classifying diverse samples. A robust machine learning model could be implemented as part of a non-invasive and low-cost aid in diagnosing SRN and other diseases.

## Materials and Methods

**Dataset.** Raw 16S rRNA gene sequence data isolated from human stool samples was downloaded from NCBI Sequence Read Archive (accession no. SRP062005) (1, 6). This dataset contains stool samples from a total of 490 subjects. For this analysis, samples from subjects identified in the metadata as normal, high risk normal, or adenoma were categorized as "normal", while samples from subjects identified as advanced adenoma or carcinoma were categorized as "screen relevant neoplasia" (SRN). The resulting dataset consisted of 261 normal samples and 229 SRN samples.

**Data Processing.** The full dataset was preprocessed with mothur (v1.47) (7) to join forward and reverse reads, merge duplicate reads, align to the SILVA reference database (v132) (8), precluster, remove chimeras with UCHIME (6), assign taxonomy, and remove non-bacterial reads following the Schloss Lab MiSeq standard operating procedure described on the mothur website (https://mothur.org/wiki/miseq_sop/). 100 splits of the 490 samples were generated where 80% of the samples (392 samples) were randomly assigned to the training set and the remaining 20% (98 samples) were assigned to the test set. Using 100 splits of the data accounts for the variation that may be observed depending on the samples that are in the training or test sets. Each sample was in the training set an average of 80 times (standard deviation = 4.1) and the

5

test set an average of 20 times (standard deviation = 4.1).

The data was processed through two workflows. In the control workflow, all the data was clustered together with OptiClust (5) to generate OTUs and the resulting abundance tables were split into the training and testing sets. In the experimental workflow, the preprocessed data was split into the training and testing sets. The training set was clustered into OTUs using OptiClust, then the test set was fit to the OTUs of the training set using the OptiFit algorithm (4). The OptiFit algorithm was run with method open so that any sequences that did not map to the existing OTU clusters would form new OTUs. For both pathways, the shared files were sub-sampled to 10,000 reads per sample.

***Machine Learning.*** A random forest model was trained with the R package mikrompl (v 1.2.0) (9) to predict the diagnosis (SRN or normal) for the samples in the test set for each data split. The training set was preprocessed to normalize OTU counts (scale and center), collapse correlated OTUs, and remove OTUs with zero variance. The preprocessing from the training set was then applied to the test set. Any OTUs in the test set that were not in the training set were removed. P values comparing model performance were calculated as previously described (10). The averaged ROC curves were plotted by taking the average and standard deviation of the sensitivity at each specificity value.

***Code Availability.*** The analysis workflow was implemented in Snakemake (11) . Scripts for analysis were written in R (12) and GNU bash (13). The software used includes mothur v1.47.0 (7), RStudio (14), the Tidyverse metapackage (15), R Markdown (16), the SRA toolkit (17), and conda (18). The complete workflow and supporting files required to reproduce this study are available at: https://github.com/SchlossLab/Armour_OptiFitGLNE_XXXX_2021

## Acknowledgements

## References

1. **Baxter NT**, **Ruffin MT**, **Rogers MAM**, **Schloss PD**. 2016. Microbiota-based model improves the sensitivity of fecal immunochemical test for detecting colonic lesions. Genome Medicine **8**:37. doi:10.1186/s13073-016-0290-3.

2. **Zackular JP**, **Rogers MAM**, **Ruffin MT,I**, **Schloss PD**. 2014. The human gut microbiome as a screening tool for colorectal cancer. Cancer Prevention Research **7**:1112–1121. doi:10.1158/1940-6207.CAPR-14-0129.

3. **Westcott SL**, **Schloss PD**. 2015. De novo clustering methods outperform reference-based methods for assigning 16S rRNA gene sequences to operational taxonomic units. PeerJ **3**:e1487. doi:10.7717/peerj.1487.

4. **Sovacool KL**, **Westcott SL**, **Mumphrey MB**, **Dotson GA**, **Schloss PD**. 2022. OptiFit: An improved method for fitting amplicon sequences to existing OTUs. mSphere **7**:e00916–21. doi:10.1128/msphere.00916-21.

5. **Westcott SL**, **Schloss PD**. 2017. OptiClust, an improved method for assigning amplicon-based sequence data to operational taxonomic units. mSphere **2**:e00073–17. doi:10.1128/mSphereDirect.00073-17.

6. **Edgar RC**, **Haas BJ**, **Clemente JC**, **Quince C**, **Knight R**. 2011. UCHIME improves sensitivity and speed of chimera detection. Bioinformatics **27**:2194–2200. doi:10.1093/bioinformatics/btr381.

7. **Schloss PD**, **Westcott SL**, **Ryabin T**, **Hall JR**, **Hartmann M**, **Hollister EB**, **Lesniewski RA**, **Oakley BB**, **Parks DH**, **Robinson CJ**, **Sahl JW**, **Stres B**, **Thallinger GG**, **Van Horn DJ**, **Weber CF**. 2009. Introducing mothur: Open-source, platform-independent, community-supported software for describing and comparing microbial communities. Applied and Environmental Microbiology **75**:7537–7541. doi:10.1128/AEM.01541-09.

8. **Quast C**, **Pruesse E**, **Yilmaz P**, **Gerken J**, **Schweer T**, **Yarza P**, **Peplies J**, **Glöckner FO**. 2013. The SILVA ribosomal RNA gene database project: Improved data processing and web-based tools. Nucleic Acids Research **41**:D590–D596. doi:10.1093/nar/gks1219.

9.    **Topçuoğlu BD**, **Lapp Z**, **Sovacool KL**, **Snitkin E**, **Wiens J**, **Schloss PD**. 2021.    mikropml: User-Friendly R Package for Supervised Machine Learning Pipelines.    Journal of Open Source Software **6**:3073. doi:10.21105/joss.03073.

10.   **Topçuoğlu BD**, **Lesniak NA**, **Ruffin MT**, **Wiens J**, **Schloss PD**. 2020. A framework for effective application of machine learning to microbiome-based classification problems. mBio **11**:e00434–20. doi:10.1128/mBio.00434-20.

11.   **Koster J**, **Rahmann S**. 2012. Snakemake–a scalable bioinformatics workflow engine. Bioinformatics **28**:2520–2522. doi:10.1093/bioinformatics/bts480.

12.   **R Core Team**. 2020. R: A language and environment for statistical computing. R Foundation for Statistical Computing, Vienna, Austria.

13.   **GNU Project**. Bash reference manual.

14.   **RStudio Team**. 2019. RStudio: Integrated development environment for r. RStudio, Inc., Boston, MA.

15.   **Wickham H**, **Averick M**, **Bryan J**, **Chang W**, **McGowan LD**, **François R**, **Grolemund G**, **Hayes A**, **Henry L**, **Hester J**, **Kuhn M**, **Pedersen TL**, **Miller E**, **Bache SM**, **Müller K**, **Ooms J**, **Robinson D**, **Seidel DP**, **Spinu V**, **Takahashi K**, **Vaughan D**, **Wilke C**, **Woo K**, **Yutani H**. 2019. Welcome to the Tidyverse. Journal of Open Source Software **4**:1686. doi:10.21105/joss.01686.

16.   **Xie Y**, **Allaire JJ**, **Grolemund G**. 2018. R Markdown: The Definitive Guide. Taylor & Francis, CRC Press.

17.   SRA-Tools - NCBI. http://ncbi.github.io/sra-tools/.

18.   2016. Anaconda Software Distribution. Anaconda Documentation. Anaconda Inc.
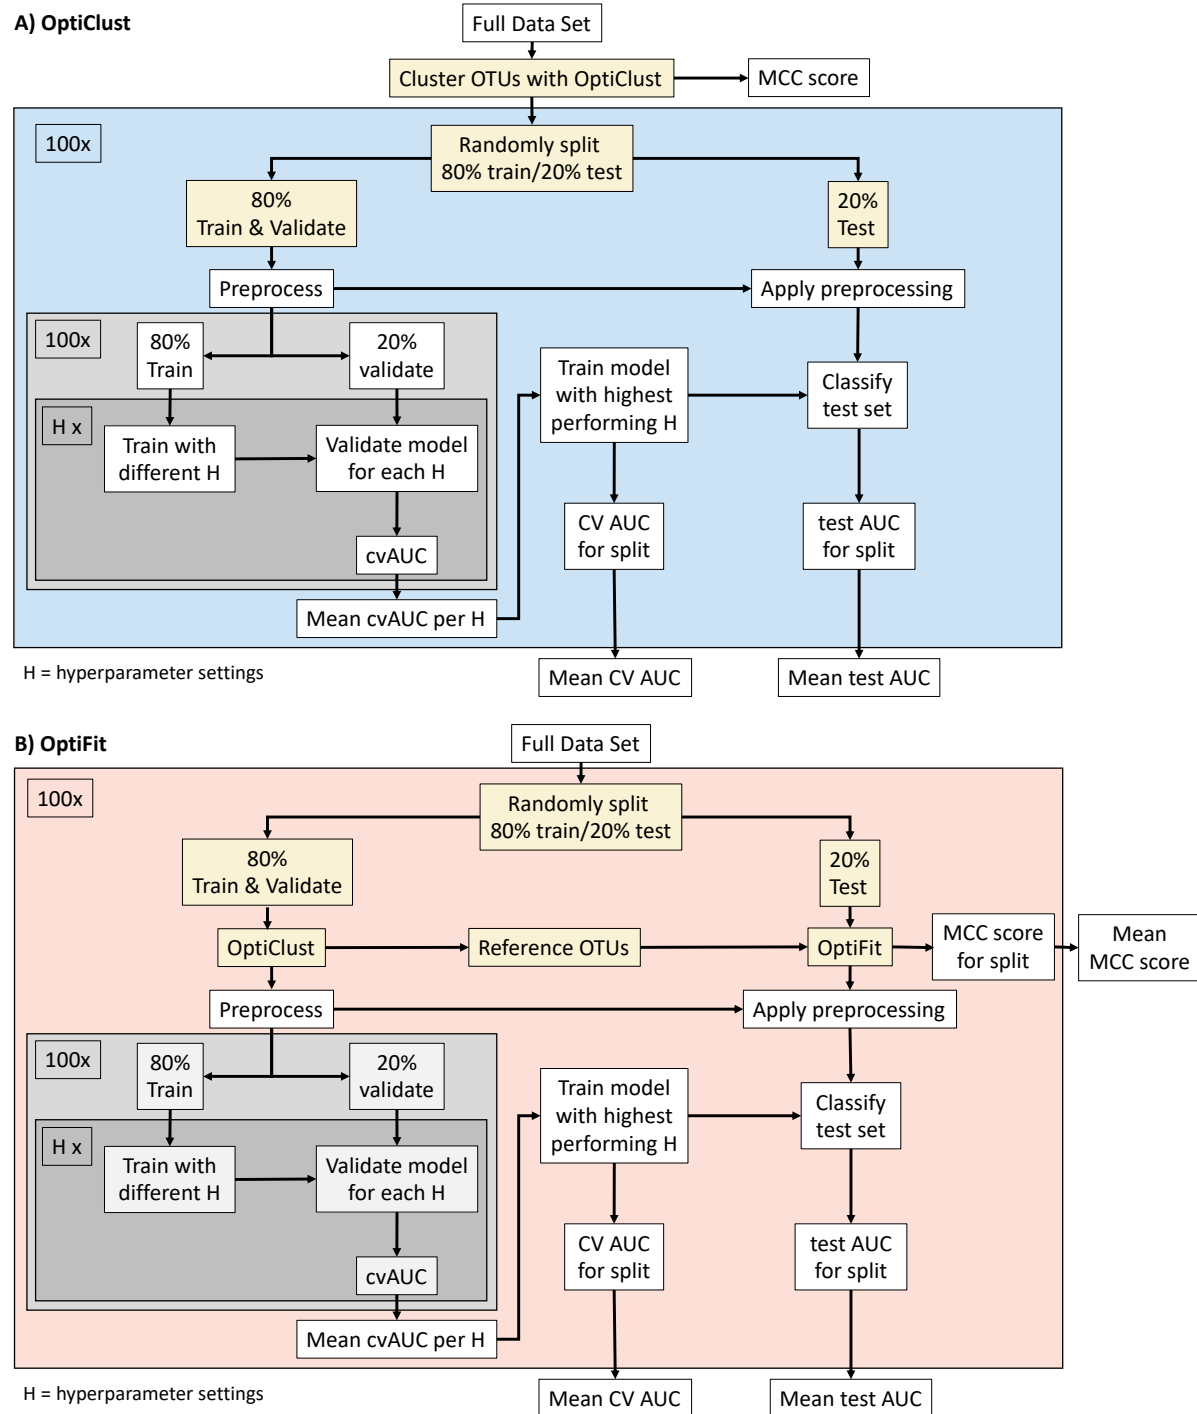
8

# Figures

**A) OptiClust**

**B) OptiFit**

**Figure 1: Workflows. A)** OptiClust workflow: The full dataset was clustered into OTUs using the OptiClust algorithm in mothur. The data was then split into two sets where 80% of the samples were assigned to the training set and 20% to the testing set. The training set was preprocessed with mikropml to normalize values (scale and center), collapse correlated features, and remove features with zero variance. Using mikropml,

9

the training set was split into train and validate sets to compare results using different hyperparameter settings. The highest performing hyperparameter setting was then used to train the model with the full training set. The preprocessing scale from the training set was applied to the test dataset, then the trained model was used to classify the samples in the test set. Based on the actual classification and predicted classification, the area under the receiver operating characteristic curve (AUROC) was calculated to summarize model performance. The entire process was repeated 100 times to account for variability depending on the split of the data resulting in a total of 100 AUROC values summarizing the performance of the standard OptiClust workflow. **B)** OptiFit workflow: The dataset was first split into two sets where 80% of the samples were assigned to the training set and 20% to the testing set. The training set was then clustered into OTUs using the OptiClust algorithm in mothur. The resulting abundance data was preprocessed with mikropml to normalize values (scale/center), collapse correlated features, and remove features with zero-variance. Using mikropml, the training set was split into train and validate sets to compare results using different hyperparameter settings. The highest performing hyperparameter setting was then used to train the model with the full training set. The OptiFit algorithm in mothur was used to cluster the held out testing dataset using the OTUs of the training set as a reference. The preprocessing scale from the training set was applied to the test dataset, then the trained model was used to classify the samples in the test set. Based on the actual classification and predicted classification, the area under the receiver operating characteristic curve (AUROC) was calculated to summarize model performance. The entire process was repeated 100 times to account for variability depending on the split of the data resulting in a total of 100 AUROC values summarizing the performance of the new OptiFit workflow.
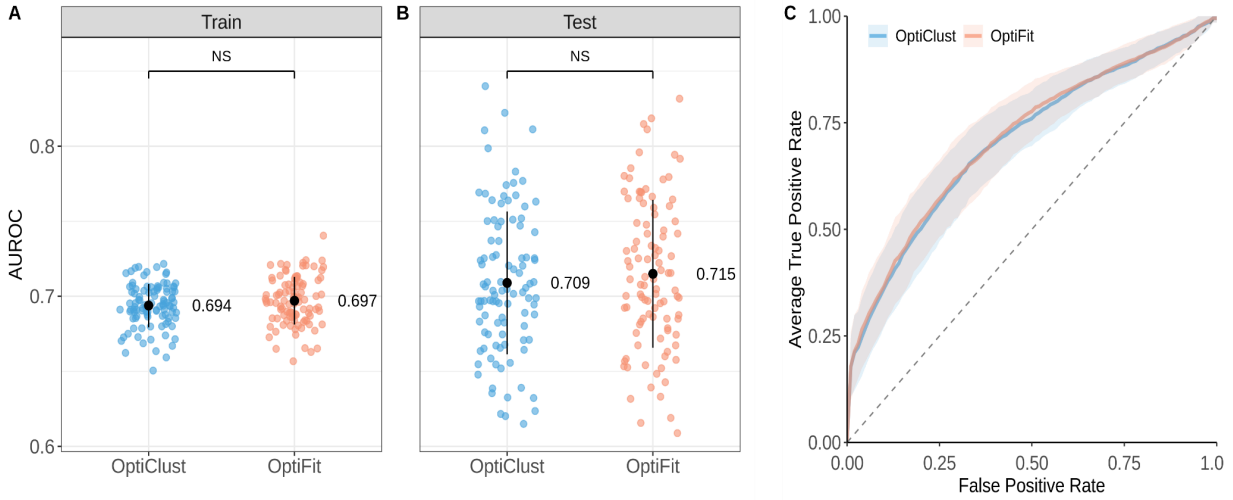
**Figure 2: Model Performance. A)** Area under the receiver operating characteristic (AUROC) curve during cross-validation for the OptiClust and OptiFit workflows. Mean and standard deviation of the AUROC is represented by the black dot and whiskers. Mean AUROC is printed to the right of the points. **B)** AUROC on the test data for the OptiClust and OptiFit workflows. Mean and standard deviation of the AUROC is represented by the black dot and whiskers. The mean AUROC is printed to the right of the points. **C)** Averaged receiver operating characteristic (ROC) curves.